

Onir Narahari

424-299-0453 | onir.narahari@utexas.edu | github.com/onir-narahari

EDUCATION

University of Texas at Austin

Bachelor of Science in Statistics and Data Science, Minor in Computer Science

– GPA: 3.7

Austin, TX

May 2028

EXPERIENCE

iNymbus

AI Developer

Remote

May 2026 – Present

- Developed an AI chatbot using **Claude** and **AWS Bedrock** to let customers query databases in plain English, replacing manual reporting for **50+ customer databases** in production
- Built a SQL guardrail layer using AST-based query validation and single-schema read-only roles, blocking cross-tenant access with **zero leaks** across **100+ security tests**
- Improved answer accuracy from **65% to 90%** by routing queries between **Claude Haiku** and **Sonnet** with self-correcting retries and semantic schema retrieval

RepoMax – tryrepomax.com

Founder

Austin, TX

Jun 2026 – Present

- Built AI GitHub analysis platform generating repo scores and resume bullets for CS students, reaching **500+ users** and **250+ repos submitted**
- Engineered RAG pipeline using GitHub APIs and embeddings to extract technical depth from repos, cutting token usage **20%** via prompt caching
- Built outreach automation using Next.js, Supabase, and Hermes to find and verify founder emails from GitHub repo context, keeping bounce rate under **5%** across **100+ sends**

Jump Finance

Software Engineer Intern

Austin, TX

Feb 2025 – Aug 2025

- Led team of 7 developers to ship full-stack course registration platform serving 100+ students.
- Built scalable infrastructure including PostgreSQL schema design, REST APIs, authentication, and React frontend supporting 250+ concurrent users with 99% uptime.
- Deployed platform on Google Cloud Platform, accelerating production launch timeline by 3 weeks.

PROJECTS

Undergraduate Researcher | ICML Submitted

March 2026 – June 2026

- Co-authored 5-page manuscript on LLM steganography, studying whether hidden messages can be embedded in model-generated text while preserving natural-sounding output.
- Designed Synchronized Logit Steering algorithm that encodes payload bits through high-entropy token-rank selection, using a 40-token generated-output window to synchronize sender and receiver decoding.
- Evaluated 4 encoding variants on gemma-2-2b across ShareGPT, GSM8K, and SWE-bench Verified, achieving 0.303 bits/token and 15x higher payload capacity while keeping most K-S detectability scores below 0.05.

MarketTrace: Real-Time Exchange Simulator

April 2026 – July 2026

- Built asynchronous electronic exchange simulator processing 100k+ orders with price-time priority matching, Level 2 order book depth, partial fills, cancellations, trade tape events, and live bid/ask updates.
- Reduced 100k-order benchmark runtime by 44x from 1,022s to 23s by redesigning the limit order book with FIFO price levels, O(1) cancellations, bounded event history, and serialization suppression.
- Implemented Avellaneda-Stoikov market maker with reservation-price skewing, volatility-aware spreads, position limits, and dynamic quote sizing, generating 5,547 dollars simulated PnL.
- Added order-flow imbalance and VPIN-style toxic-flow signals to adjust spreads, cancel stale quotes, and reduce inventory exposure under asymmetric liquidity and adverse-selection risk.
- Built FastAPI WebSocket market data server and React dashboard streaming sequenced snapshots, batched trades, spread, depth, volume, fill rate, average spread, inventory, PnL, and agent metrics.

TECHNICAL SKILLS

Languages: Python, Java, JavaScript, SQL, HTML/CSS

AI/ML: PyTorch, scikit-learn, XGBoost, Random Forest, Pandas, NumPy, LangChain, LlamaIndex, Retrieval-Augmented Generation (RAG), Vector Embeddings, OCR (Tesseract, PaddleOCR), Knowledge Distillation

Backend & APIs: FastAPI, Flask, REST APIs, OpenAI API

Databases: PostgreSQL, FAISS, ChromaDB

Cloud & Tools: Google Cloud Platform, Git, Docker, Linux, Cron Jobs